

Heart Disease Prediction

From the raw UCI Cleveland dataset to a **deployed, monitored prediction API** — covering experiment tracking, testing, CI/CD, containerization, Kubernetes and observability.

0.960

TEST ROC-AUC
LOGISTIC RE-
GRESSION

303

PATIENTS
13 CLINICAL
FEATURES

24

PYTEST TESTS
IN 4-STAGE CI

2

K8S REPLICAS
BEHIND LOADBAL-
ANCER

REPOSITORY	github.com/rajsingha/heart-disease-mlops
DATASET	UCI Heart Disease — Cleveland subset
STACK	scikit-learn · MLflow · FastAPI · Docker · Kubernetes · Prometheus · Grafana
DATE	July 2026

CONTENTS

01	Project overview	2	07	Packaging & reproducibility	5
02	Dataset	2	08	CI/CD pipeline	6
03	EDA findings	3	09	Containerization	6
04	Preprocessing & feature engineering	4	10	Kubernetes deployment	7
05	Model development & comparison	4	11	Monitoring & logging	7
06	Experiment tracking with MLflow	5	12	Setup instructions	8


SECTION 01

Project overview

The goal of this project was to take a machine learning model from raw data all the way to a deployed, monitored API using standard MLOps practices. I used the UCI Heart Disease dataset to train a binary classifier that predicts whether a patient has heart disease, and then built everything around it: experiment tracking with MLflow, unit tests, a CI pipeline on GitHub Actions, a Docker image, a Kubernetes deployment, and basic monitoring with Prometheus and Grafana.

The final model is a tuned logistic regression with a test ROC-AUC of **0.96**. It is served by a FastAPI app with a `/predict` endpoint that takes patient data as JSON and returns the prediction along with a probability.

Architecture



SECTION 02

Dataset

I used the Cleveland subset of the UCI Heart Disease dataset: **303 patients** and **13 clinical features**. The original label (`num`) goes from 0 to 4, so I converted it to a binary target where 1 means any presence of heart disease.

FEATURE	DESCRIPTION	TYPE
age	Age in years	numeric
sex	1 = male, 0 = female	binary
cp	Chest pain type (1–4, 4 = asymptomatic)	categorical
trestbps	Resting blood pressure (mm Hg)	numeric
chol	Serum cholesterol (mg/dl)	numeric
fbs	Fasting blood sugar > 120 mg/dl	binary
restecg	Resting ECG result (0–2)	categorical
thalach	Maximum heart rate achieved	numeric
exang	Exercise-induced angina	binary
oldpeak	ST depression (exercise vs rest)	numeric
slope	Slope of peak exercise ST segment (1–3)	categorical
ca	Major vessels colored by fluoroscopy (0–3)	numeric
thal	Thalassemia (3/6/7)	categorical

The download is scripted in `src/data/download.py`. It first tries the `ucimlrepo` package and falls back to the raw file on the UCI archive if that fails, so the raw data can always be rebuilt from a fresh checkout. I also committed the cleaned CSV under `data/processed/` so the CI pipeline can retrain the model without depending on the UCI servers being up.

SECTION 03

EDA findings

The figures are generated by `python -m src.eda` into `reports/figures/`, and the same analysis can be run interactively in `notebooks/eda.ipynb`. The main things I found:

- 01 The classes are close to balanced: 164 patients without disease (54.1%) and 139 with disease (45.9%). Because of this I didn't do any resampling and just used stratified splits. →

`class_distribution.png`

- 02 Missing values are rare. Only `ca` (4 rows) and `thal` (2 rows) have them, written as `?` in the raw file. I impute them inside the model pipeline (median for numeric, most frequent for categorical) so the imputers are fitted only on training data and there is no leakage. → `missing_values.png`

- 03 Some features separate the classes clearly. Patients with disease reach lower maximum heart rates, show more ST depression during exercise, and have more vessels colored during fluoroscopy. → [histograms_numeric.png](#)
- 04 On the categorical side, asymptomatic chest pain (`cp = 4`), exercise-induced angina, a flat or down-sloping ST segment, and a reversible thalassemia defect (`thal = 7`) all come with much higher disease rates. → [categorical_vs_target.png](#)
- 05 The correlation heatmap showed no pair of features above roughly 0.6, so I kept all of them. → [correlation_heatmap.png](#)

SECTION 04

Preprocessing and feature engineering

All preprocessing lives in a single sklearn `ColumnTransformer` (`src/data/preprocess.py`) that becomes the first step of every model pipeline:

- **Numeric features** (`age, trestbps, chol, thalach, oldpeak, ca`): median imputation, then standard scaling.
- **Categorical features** (`sex, cp, fbs, restecg, exang, slope, thal`): most frequent imputation, then one-hot encoding with `handle_unknown="ignore"`.

DESIGN CHOICE

The transformer is saved inside the pipeline itself. The API loads one joblib file and can feed raw JSON values straight into it. There is no second preprocessing implementation that could drift away from the training code.

The data is split 80/20 with stratification and a fixed random seed (42). All hyperparameter tuning happens with cross-validation inside the training split, and the test set is only used at the end for the final comparison.

SECTION 05

Model development and comparison

I tuned three model families with `GridSearchCV`, using stratified 5-fold cross-validation and refitting on ROC-AUC (`src/models/train.py`):

FAMILY	GRID	BEST PARAMS
Logistic Regression	$C \in \{0.01, 0.1, 1, 10\}$, L1/L2 penalty	$C = 1.0$, L2
Random Forest	200/400 trees, depth None/5/10, min_split 2/5	200 trees, depth 5, min_split 5
XGBoost	200/400 trees, depth 3/5, lr 0.05/0.1	200 trees, depth 3, lr 0.05

Results on the held-out test set (61 patients, full table in `reports/model_comparison.csv`):

MODEL	CV ROC-AUC	TEST ROC-AUC	ACCURACY	PRECISION	RECALL	F1
Logistic Regression ✓	0.908 ± 0.018	0.960	0.869	0.813	0.929	0.867
Random Forest	0.901 ± 0.025	0.949	0.885	0.839	0.929	0.881
XGBoost	0.865 ± 0.029	0.944	0.869	0.813	0.929	0.867

I picked **logistic regression**. It has the best cross-validated and test ROC-AUC and the lowest variance across folds. With only 303 rows the tree ensembles don't get to use their extra capacity, and the simpler linear model is less likely to overfit. Recall is high (0.93), which matters for a screening scenario where missing a sick patient is worse than a false alarm. The coefficients are also easy to interpret, which is a plus in a medical setting.

Confusion matrices and ROC curves for all three models are saved in `reports/figures/` and attached to the MLflow runs.

SECTION 06

Experiment tracking with MLflow

Each model family gets its own MLflow run (local sqlite store `mlflow.db`, artifacts under `mlruns/`). For every run I log:

- the best hyperparameters and the CV configuration
- cross-validation means for accuracy, precision, recall, F1 and ROC-AUC, plus the ROC-AUC standard deviation
- all test metrics
- the confusion matrix and ROC curve as PNGs, a text classification report, and the fitted pipeline itself via `mlflow.sklearn.log_model`

The UI can be started with `mlflow ui --backend-store-uri sqlite:///mlflow.db`. Screenshots are in `screenshots/mlflow_runs.png` and `screenshots/mlflow_run_detail.png`.

Next to the MLflow store, the winning pipeline is exported to `models/model.joblib` together with `models/model_metadata.json`, which records the metrics, hyperparameters, library versions, training timestamp and the MLflow run id. The API reads this file and reports the model family in its `/health` response.

SECTION 07

Packaging and reproducibility

- `requirements.txt` pins the full training environment.

- `requirements-api.txt` pins only what the serving container needs. The scikit-learn and xgboost versions match the training environment exactly so the pickled pipeline loads identically.
- The preprocessing and the model are one artifact, so there is no way to serve the model with the wrong preprocessing.
- The cleaned dataset is committed and every random operation uses seed 42, so training is reproducible. CI proves this by retraining from scratch on every push.
- Each stage is runnable as a module (`python -m src.data.download` , `python -m src.models.train` , and so on).

SECTION 08

CI/CD pipeline

The GitHub Actions workflow (`.github/workflows/ci.yml`) runs on every push and pull request to `main` . It has four jobs that run in sequence, and a failure in any job stops the rest:

JOB	WHAT IT DOES
Lint	<code>ruff check src tests</code>
Test	runs the 24 pytest tests (data cleaning, pipeline, training, MLflow logging, API)
Train	retrains from the committed cleaned CSV and uploads the model, MLflow store and reports as artifacts
Docker	builds the image with the fresh model, starts the container, waits for <code>/health</code> , and POSTs a sample patient to <code>/predict</code>

The tests run against synthetic data generated in `tests/conftest.py` , so they don't need network access or a pre-trained model. A green run is shown in `screenshots/ci_pipeline.png` .

SECTION 09

Containerization

The `Dockerfile` starts from `python:3.12-slim` , installs the pinned serving dependencies, copies `src/` and the trained model, and runs uvicorn as a non-root user. A `HEALTHCHECK` hits `/health` every 30 seconds.

I verified it locally:

LOCAL VERIFICATION

```

docker build -t heart-disease-api:latest .
docker run --rm -p 8000:8000 heart-disease-api:latest

curl http://localhost:8000/health
{"status":"ok","model_loaded":true,"model_family":"logistic_regression"}

curl -X POST http://localhost:8000/predict \
  -H "Content-Type: application/json" -d @patient.json
{"prediction":1,"label":"heart_disease","probability":0.8986,
  "model_family":"logistic_regression"}

```

The container returns exactly the same prediction as the local environment, which was the point of the exercise. See [screenshots/docker_run.png](#) and [screenshots/predict_response.png](#).

SECTION 10

Kubernetes deployment

`k8s/deployment.yaml` runs 2 replicas with CPU/memory requests and limits, and readiness/liveness probes on `/health`. `k8s/service.yaml` exposes them through a LoadBalancer on port 80. I deployed on Docker Desktop Kubernetes, where the LoadBalancer binds to localhost; on Minikube you would use `minikube service heart-disease-api --url` instead, and on a cloud cluster it becomes a real load balancer.

DEPLOY & VERIFY

```

kubectl apply -f k8s/deployment.yaml -f k8s/service.yaml
kubectl get pods,svc -l app=heart-disease-api

```

Both pods came up healthy and `/predict` works through the service (see [screenshots/kubectl_get_pods.png](#) and [screenshots/k8s_predict.png](#)). The full walkthrough is in [k8s/README.md](#).

SECTION 11

Monitoring and logging

Three layers of visibility:

- 01 A logging middleware in the API writes one line per request with the method, path, status code and latency. Predictions are also logged with their probability.
- 02 The `/metrics` endpoint exposes Prometheus metrics: standard HTTP request counters and latency histograms, plus a custom `model_predictions_total{predicted_class=...}` counter. Watching the class mix of predictions over time is a cheap early warning for data drift.

- 03 `docker compose up --build` starts the API together with Prometheus and Grafana. Grafana is provisioned automatically with a dashboard showing request rate, p95 latency, predictions by class and HTTP status codes (see [screenshots/grafana_dashboard.png](#)).

WHY THIS MATTERS FOR ML

For an ML service this is more than uptime monitoring. A shift in the prediction distribution or a latency regression usually shows up here long before anyone looks at model accuracy again.

SECTION 12

Setup instructions

The short version (Python 3.12+, more detail in the README):

QUICKSTART

```
python -m venv .venv
.venv\Scripts\activate          # on Linux/macOS: source .venv/bin/activate
pip install -r requirements.txt

python -m src.data.download
python -m src.data.preprocess
python -m src.eda
python -m src.models.train

pytest -v
ruff check src tests

uvicorn src.api.main:app --port 8000
```

Swagger UI is at <http://localhost:8000/docs> once the server is running.